

Robotic Navigation and Exploration

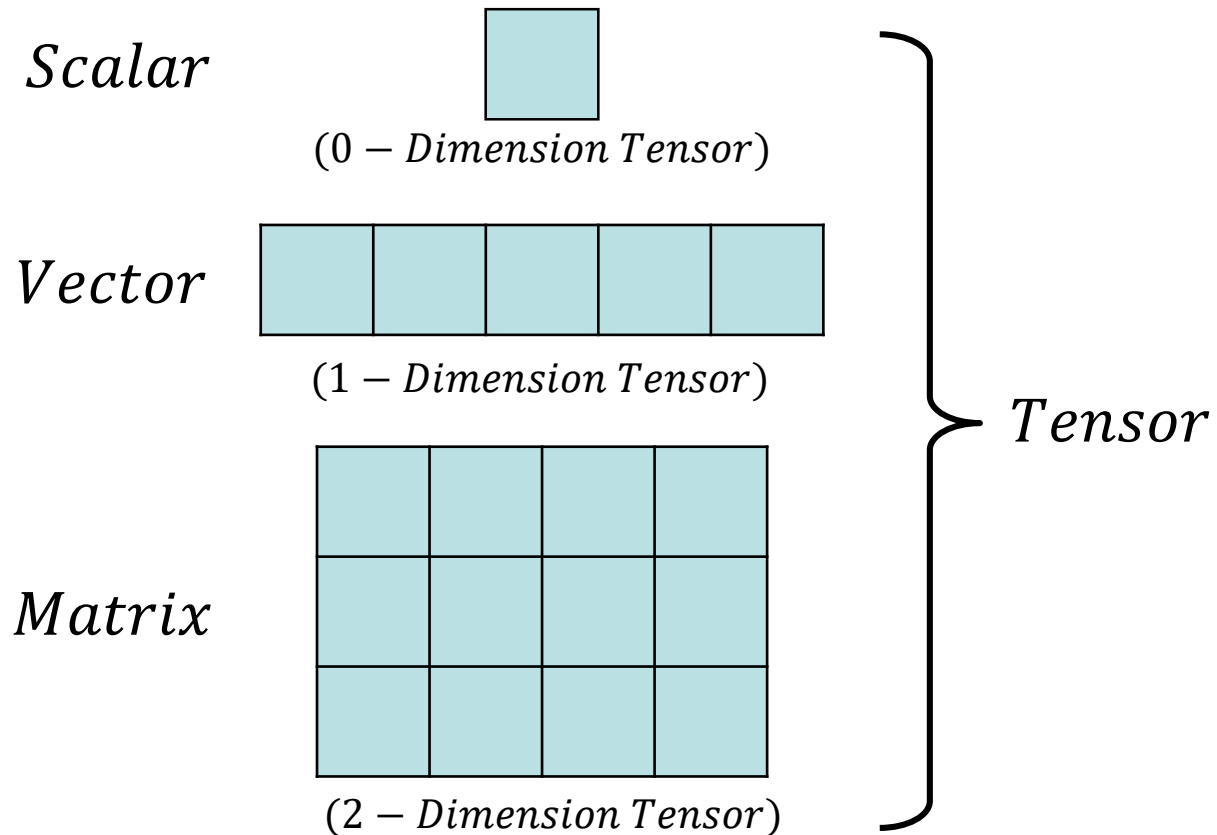
Lab5: Semantic Segmentation with PyTorch

Min-Chun Hu anitahu@cs.nthu.edu.tw
CS, NTHU

PyTorch Basics


Tensor

- In Pytorch , tensor is the basic element which can be represented as multi-dimensional array.



Computation Graph

- The core of modern deep learning library such as **Tensorflow** and **PyTorch** is the **Automatic Differentiator on Computation Graph**.
- Every neural network model can be represented as a graph. By passing the error message through the graph, the library can get the updating direction of each variables.



Computation Graph

- Linear Regression

$$\begin{bmatrix} y_1 & y_2 & y_3 & y_4 \end{bmatrix} = \begin{bmatrix} w & b \end{bmatrix} \begin{bmatrix} x_1 & x_2 & x_3 & x_4 \\ 1 & 1 & 1 & 1 \end{bmatrix}$$

```
import torch
x = torch.tensor([[0,2,3,4],[1,1,1,1]], dtype=float)
y = torch.tensor([[1,3,5,8]], dtype=float)
w = torch.tensor([[0.1,0.1]], dtype=float, requires_grad=True)

for i in range(20):
    # Forward
    y_out = torch.matmul(w,x)
    loss = (0.5 * (y - y_out)**2).sum()
    print(loss)

    # Backward
    loss.backward()
    with torch.no_grad():
        w -= 0.01 * w.grad
    w.grad.zero_()
```

Create the gradient buffer.

Every operation dynamically create the graph.

Backward operation will delete the graph.

Update parameters, prevent to building graph.

Notification

- The backward node should be single value, otherwise you need to give the weighting parameter with same size of the backward node.
 - Ex. `loss = 0.5*(y-y_out)**2`
`loss.backward(torch.tensor([[1,1,1,1]]), dtype=float))`
- Backward operation will delete the graph by default, you can set the backward parameter to retain the graph.
 - Ex. `loss.backward(retain_graph=True)`
- After backward operation, each leaf node that requires gradient will “add” the gradient at the buffer, thus you need to set the gradient buffer to zero in general usage before backward. (However, you can also use this feature to achieve special effects.)

Commonly Used Libraries

- torch.nn: Neural Network Operation.
- torch.nn.functional: Functional Operation.
- torch.optim: Optimizer.
- torch.utils.data: Dataset / Data Loader.
- torchvision: Computer vision-related operation/datasets/models.

PyTorch Documentation: <https://pytorch.org/docs/stable/index.html>

PyTorch LeNet Example

```
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim

class LeNet(nn.Module):
    def __init__(self):
        super().__init__()
        # Create Parameters Here
        self.conv1 = nn.Conv2d(1,6,kernel_size=5)
        self.conv2 = nn.Conv2d(6,16,kernel_size=5)
        self.fc1 = nn.Linear(16*5*5,128)
        self.fc2 = nn.Linear(128,10)

    def forward(self, x):
        # Construct Forward Path
        output = F.relu(self.conv1(x)) # (1,32,32) -> (6,28,28)
        output = F.max_pool2d(output) # (6,28,28) -> (6,14,14)
        output = F.relu(self.conv2(output)) # (6,14,14) -> (16,10,10)
        output = F.max_pool2d(output) # (16,10,10) -> (16,5,5)
        output = F.relu(self.fc1(output.view(-1,16*5*5))) # (16*5*5) -> (128)
        output = self.fc2(output) # (128) -> (10)
        return output
```


PyTorch LeNet Example

```
net = LeNet() # Create Network
criterion = nn.BCEWithLogitsLoss() # Loss Function
optimizer = optim.Adam(net.parameters(), lr=0.001) # Create Optimizer

for i in range(1000):
    optimizer.zero_grad() # Clear the gradient buffer
    inputs, labels = # Get data batch
    outputs = net(inputs) # Forward Propagation
    loss = criterion(outputs, labels) # Compute Loss
    loss.backward() # Backward Propagation
    optimizer.step() # Update Parameters
```

Module Block

- Inheritat “nn.Module” to define new computation block to make sure the parameters can be found recursively.

```
class SubBlock(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(16,2,kernel_size=3)
        self.conv2 = nn.Conv2d(2,1,kernel_size=3)

    def forward(self, x):
        output = self.conv1(x)
        output = self.conv2(output)
        return output
```

```
class MainNet(nn.Module):
    def __init__(self):
        super().__init__()
        # Create Parameter Here
        self.conv1 = nn.Conv2d(3,8,kernel_size=3)
        self.conv2 = nn.Conv2d(8,16,kernel_size=3)
        self.subblock = SubBlock()

    def forward(self, x):
        # Construct Forward Path
        output = self.conv1(output)
        output = F.relu(output)
        output = self.conv2(output)
        output = self.subblock(output)
        output = F.sigmoid(output)
        return output
```

Dataset and Dataloader

- `torch.utils.data.Dataset` is an abstract class representing a dataset. Your custom dataset should inherit Dataset and override the following methods:

- `__len__` : returns the size of the dataset.
- `__getitem__` : support the indexing that can be used to get ith sample.

- `torch.utils.data.DataLoader` shuffle the data in the dataset and construct an iterator for us to get a random batch of training data.


- Ex.

```
train_data = CustomDataset(...)
train_loader = DataLoader(train_data, batch_size=32, shuffle=True)

for iter, batch in enumerate(train_loader):
    inputs = torch.FloatTensor(batch['X'])
    labels = torch.FloatTensor(batch['Y'])
```

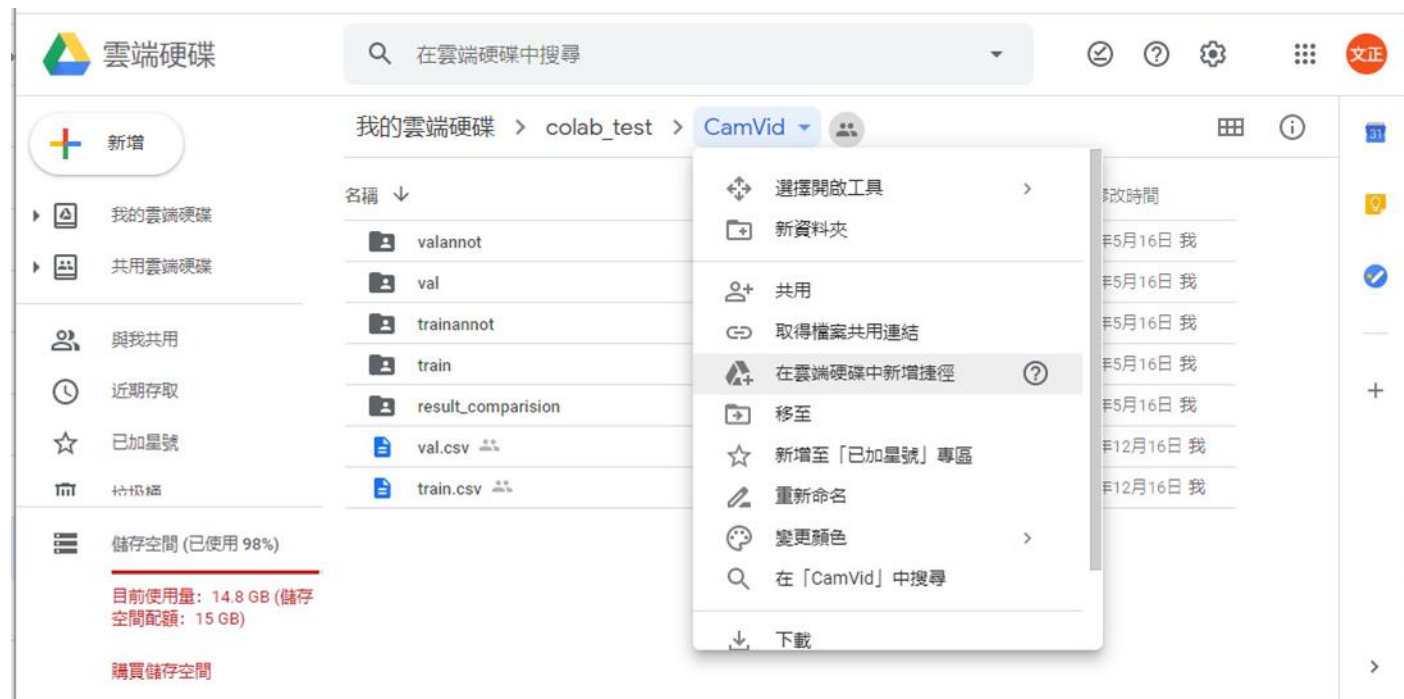
Semantic Segmentation

CamVid Dataset



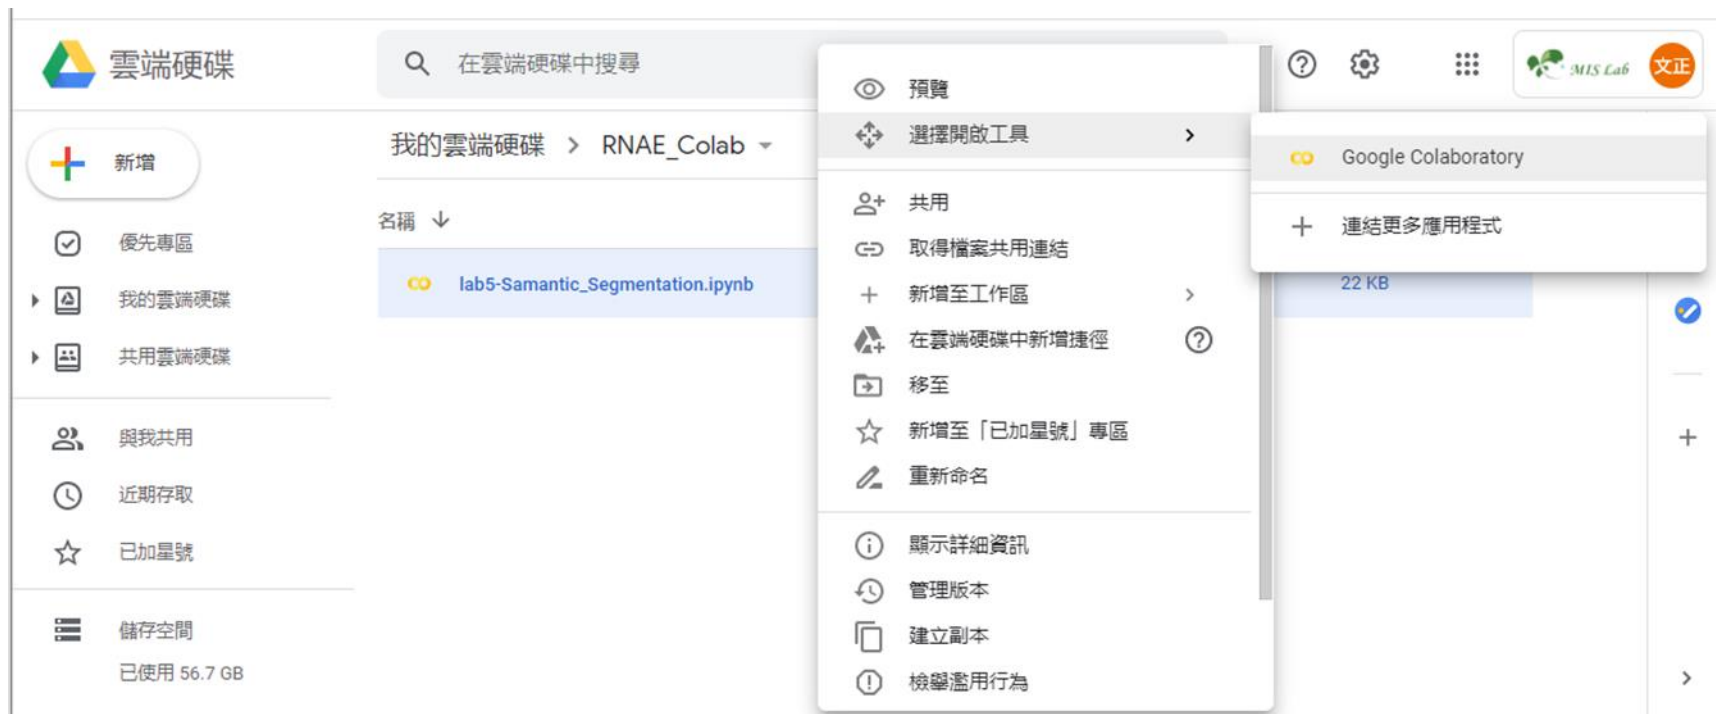
Prepare Dataset

- Link: https://drive.google.com/drive/folders/1K8kys7gWfl-As3lCz6O4Y6ETdn_OVTd?usp=sharing
- Login your google account and add the shortcut to the dataset folder to your cloud drive.



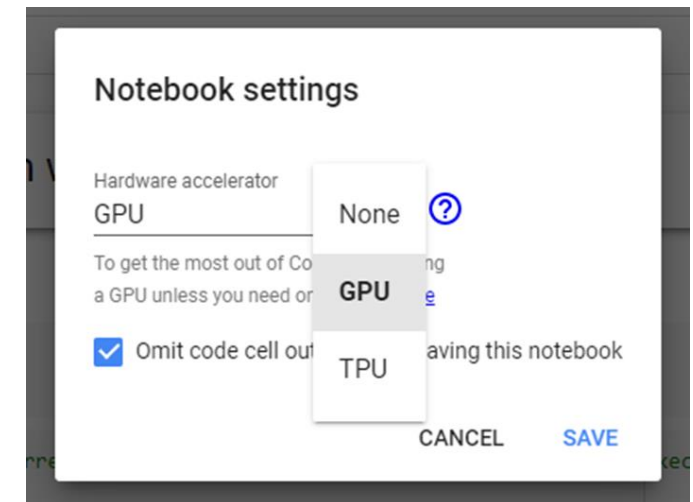

Google Colaboratory

- Upload the ipynb file to your google drive, and open it using Google Colaboratory. (If you donot have Google Colaboratory, you need to click “連結更多應用程式” and install it.)



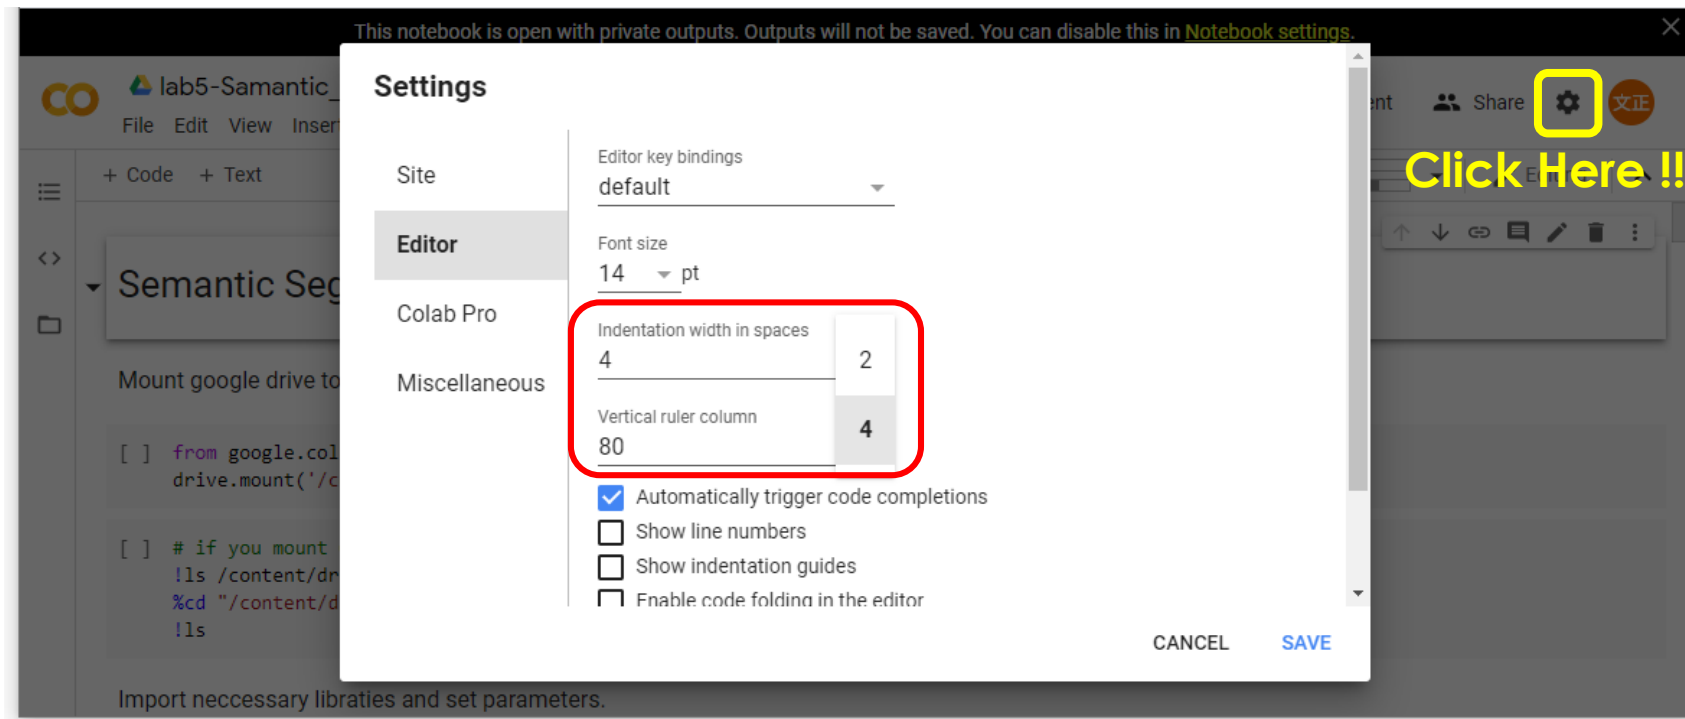
Google Colaboratory

- Set runtime hardware accelerator to “GPU”



Google Colaboratory

- Set the indentation size to 4.



Mount Google Drive

- Run the first cell and go to the URL on the message. Copy the authorization code and press “Enter”.

The screenshot shows a Google Colab notebook titled "lab5-Samantic_Segmentation.ipynb". The first code cell contains the following Python code:

```
from google.colab import drive
drive.mount('/content/drive')
```

Below the code cell, a message states: "Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3Aietf%3Aawg%3Aoauth%3A2.0%3Aoob... Enter your authorization code:".

A separate window titled "Google 登入" (Google Login) is shown, displaying the instruction: "請複製這組授權碼，然後切換至您的應用程式，再貼上授權碼：" (Please copy this authorization code, then switch to your application, and paste the authorization code:). The authorization code input field is highlighted with a red arrow pointing from the URL in the notebook.



Mount Google Drive

- Run the second cell to check the dataset file.

```
# if you mount Google drive correctly, the following commands should be able to executed correctly
!ls /content/drive/
%cd "/content/drive/My Drive/CamVid"
!ls
```

```
↳ 'My Drive' 'Shared drives'
/content/drive/.shortcut-targets-by-id/1K8kys7gWfI-As3lCz6O4Y6ETdn_OVTd_/CamVid
result_comparision  train  trainannot  train.csv  val  valannot  val.csv
```

Data Augmentation



Crop


Flip

```
# crop images and labels
w, h = img.size
if self.rand_crop:
    A_x_offset = np.int32(np.random.randint(0, w - self.new_w + 1, 1))[0]
    A_y_offset = np.int32(np.random.randint(0, h - self.new_h + 1, 1))[0]
else:
    A_x_offset = int((w - self.new_w)/2)
    A_y_offset = int((h - self.new_h)/2)

img = img.crop((A_x_offset, A_y_offset, A_x_offset + self.new_w, A_y_offset + self.new_h))
label_image = label_image.crop((A_x_offset, A_y_offset, A_x_offset + self.new_w, A_y_offset + self.new_h))

# flip images and labels
img = np.transpose(img, (2, 0, 1)) / 255.
label = np.asarray(label_image)
if np.random.sample() < self.flip_rate:
    img = np.fliplr(img)
    label = np.fliplr(label)
```

VGG-16 Feature Extractor



VGG-16 Feature Extractor

```
class Vgg16(nn.Module):
    def __init__(self, pretrained = True):
        super(Vgg16, self).__init__()
        self.vggnet = models.vgg16(pretrained)
        del(self.vggnet.classifier) # Remove fully connected layer to save memory.
        features = list(self.vggnet.features)
        self.layers = nn.ModuleList(features).eval()

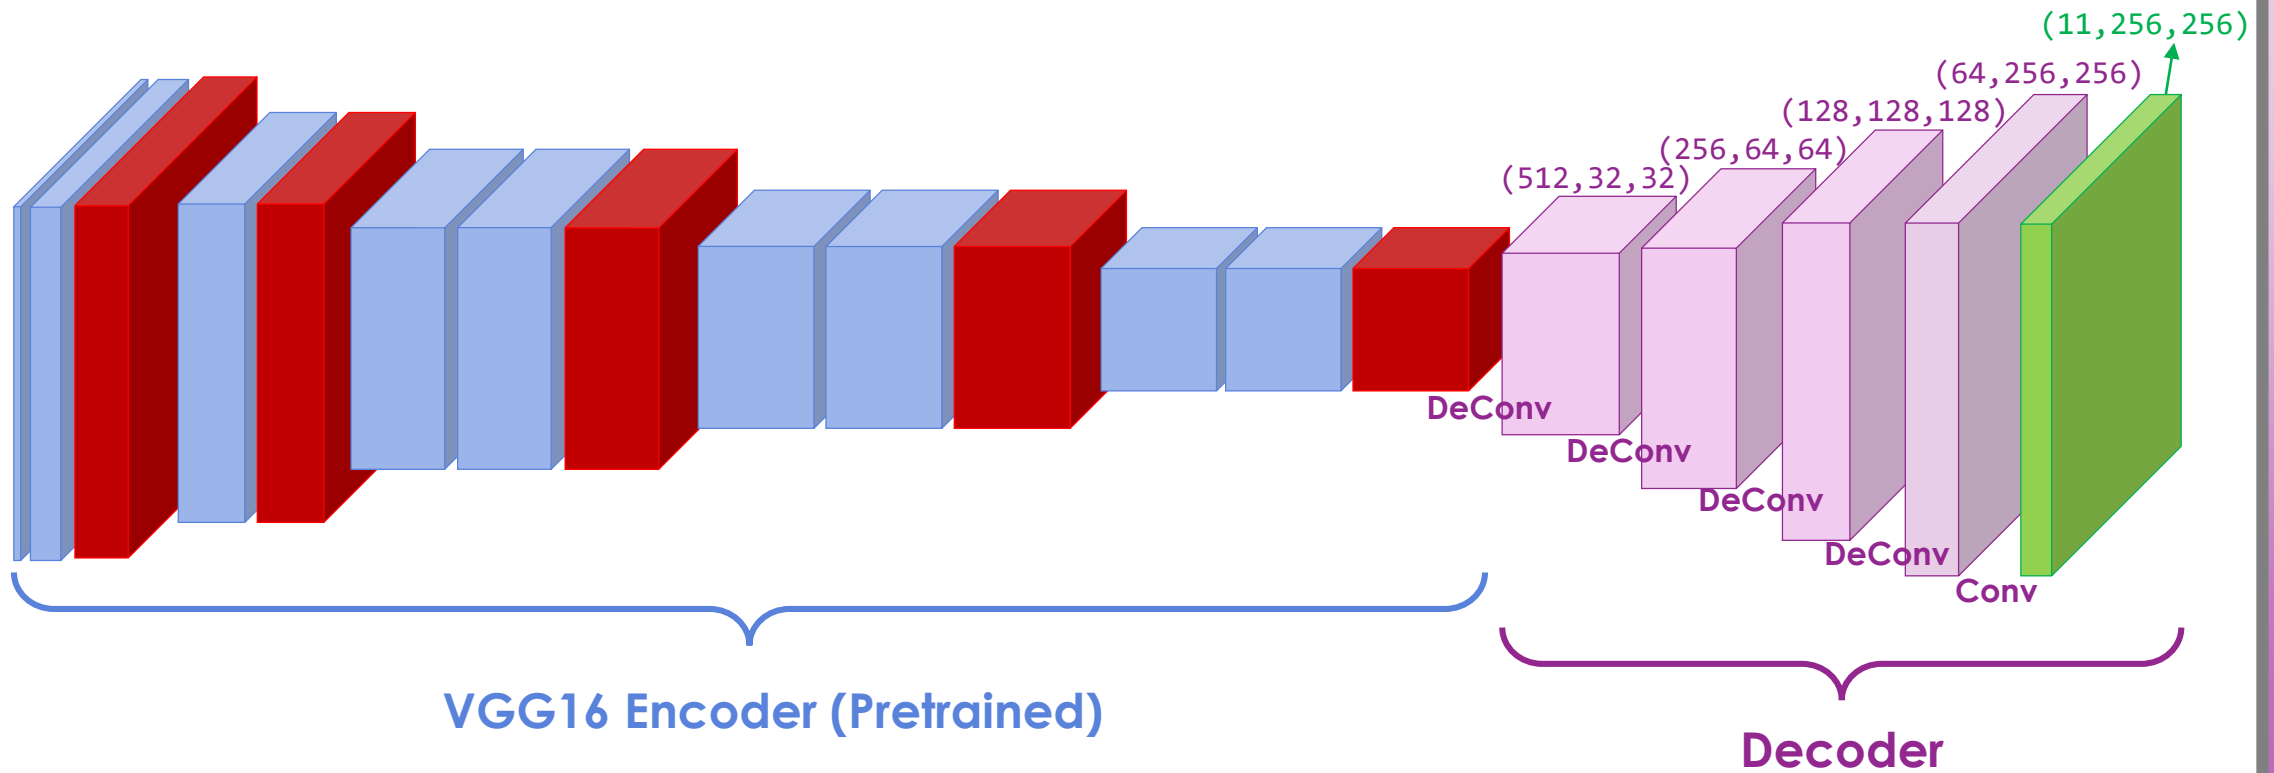
    def forward(self, x):
        results = []
        for ii,model in enumerate(self.layers):
            x = model(x)
            if ii in [3,8,15,22,29]:
                results.append(x) # (64,256,256), (128,128,128), (256,64,64), (512,32,32), (512,16,16)
        return results

vgg_model = Vgg16()
vgg_model = vgg_model.cuda()
print(vgg_model.layers)
```

VGG-16 Feature Extractor


```
ModuleList(
  (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (1): ReLU(inplace=True)
  (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (3): ReLU(inplace=True)
  (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (6): ReLU(inplace=True)
  (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (8): ReLU(inplace=True)
  (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (11): ReLU(inplace=True)
  (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (13): ReLU(inplace=True)
  (14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (15): ReLU(inplace=True)
  (16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (17): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (18): ReLU(inplace=True)
  (19): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (20): ReLU(inplace=True)
  (21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (22): ReLU(inplace=True)
  (23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (24): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (25): ReLU(inplace=True)
  (26): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (27): ReLU(inplace=True)
  (28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (29): ReLU(inplace=True)
  (30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
```

Encoder-Decoder Architecture




Up-sample Deconvolution

Standard Deconvolution




Up-sample Deconvolution



```
class DeConv2d(nn.Module):
    def __init__(self, in_channel, out_channel, kernel_size, stride, padding, dilation):
        super().__init__()
        self.up = nn.Upsample(scale_factor=2, mode='nearest')
        self.conv = nn.Conv2d(in_channel, out_channel, kernel_size=kernel_size, \
                               stride=stride, padding=padding, dilation=dilation)


    def forward(self, x):
        output = self.up(x)
        output = self.conv(output)
        return output
```

Practice - FCN

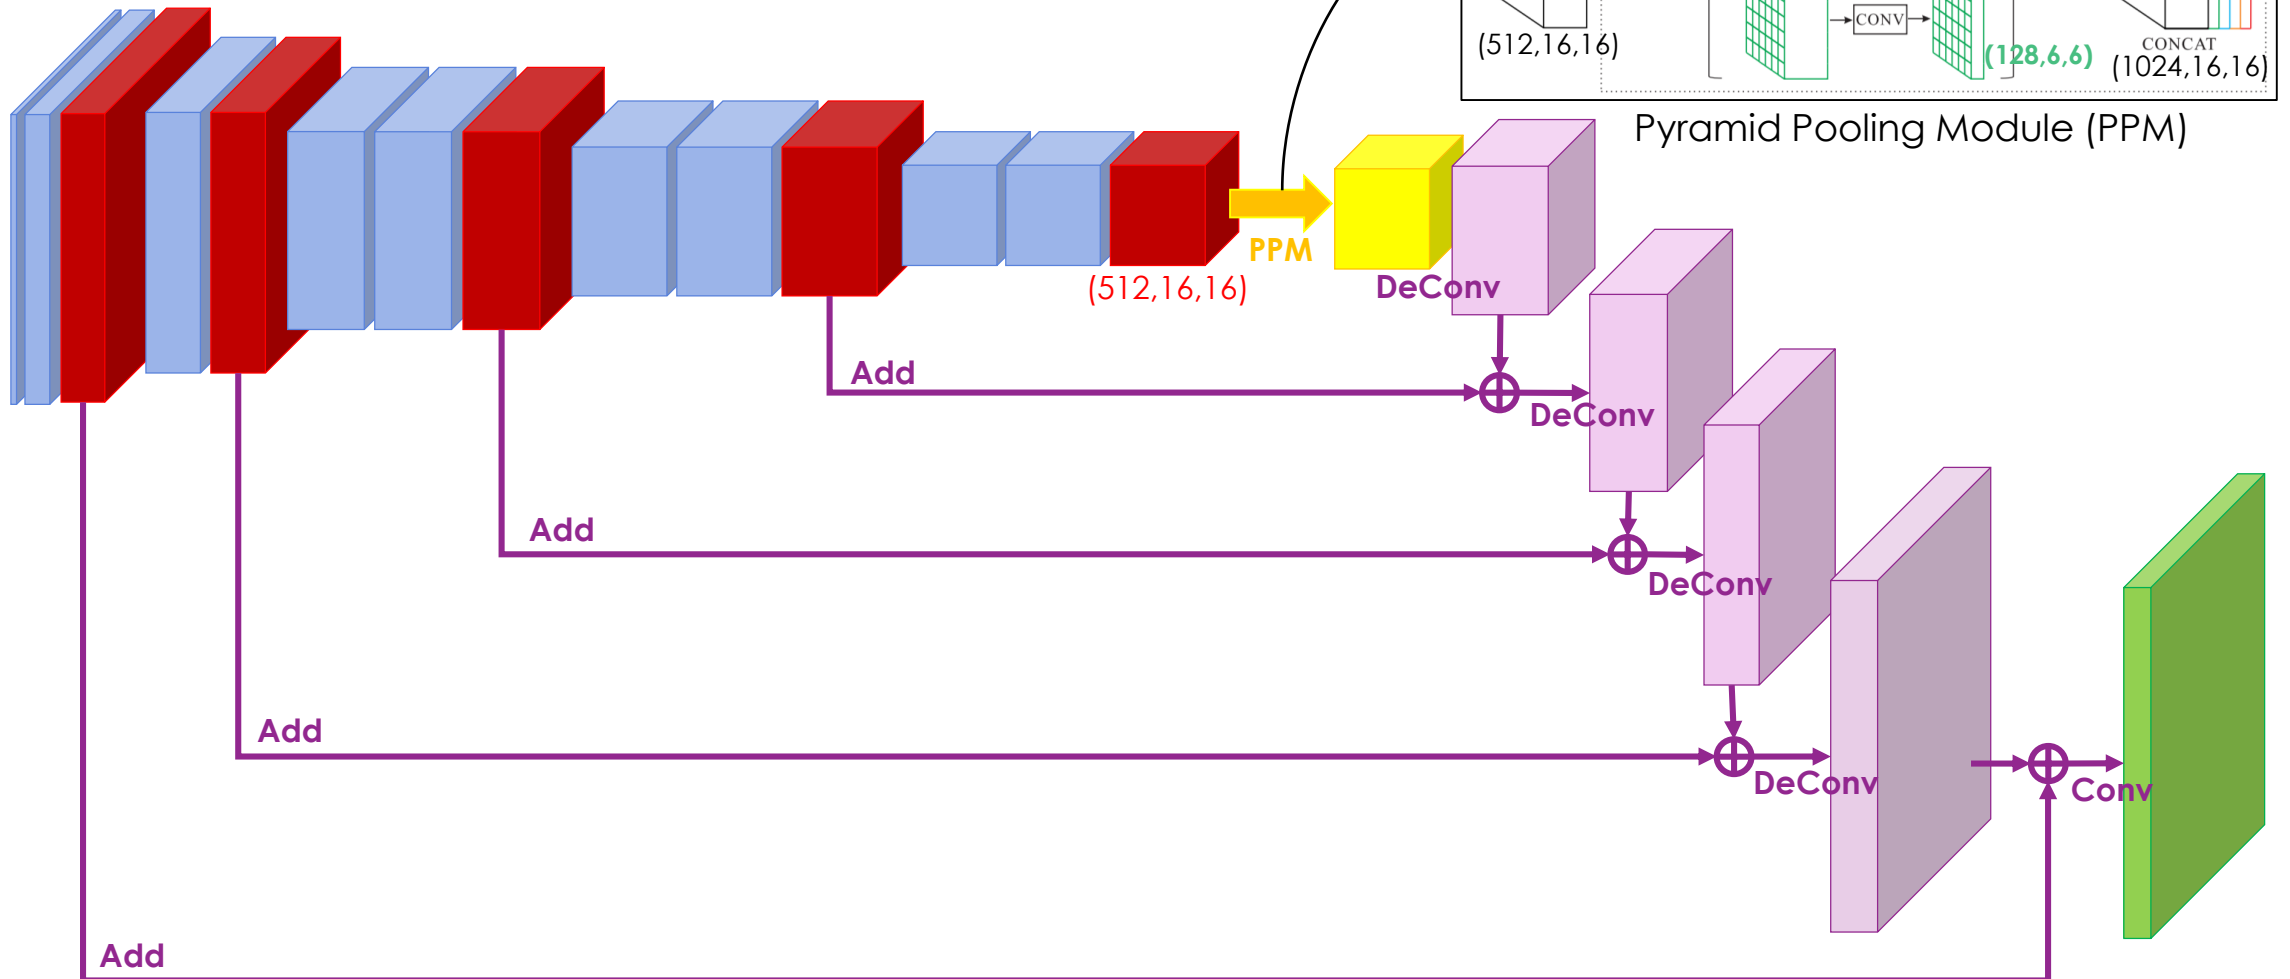


Practice - U-Net

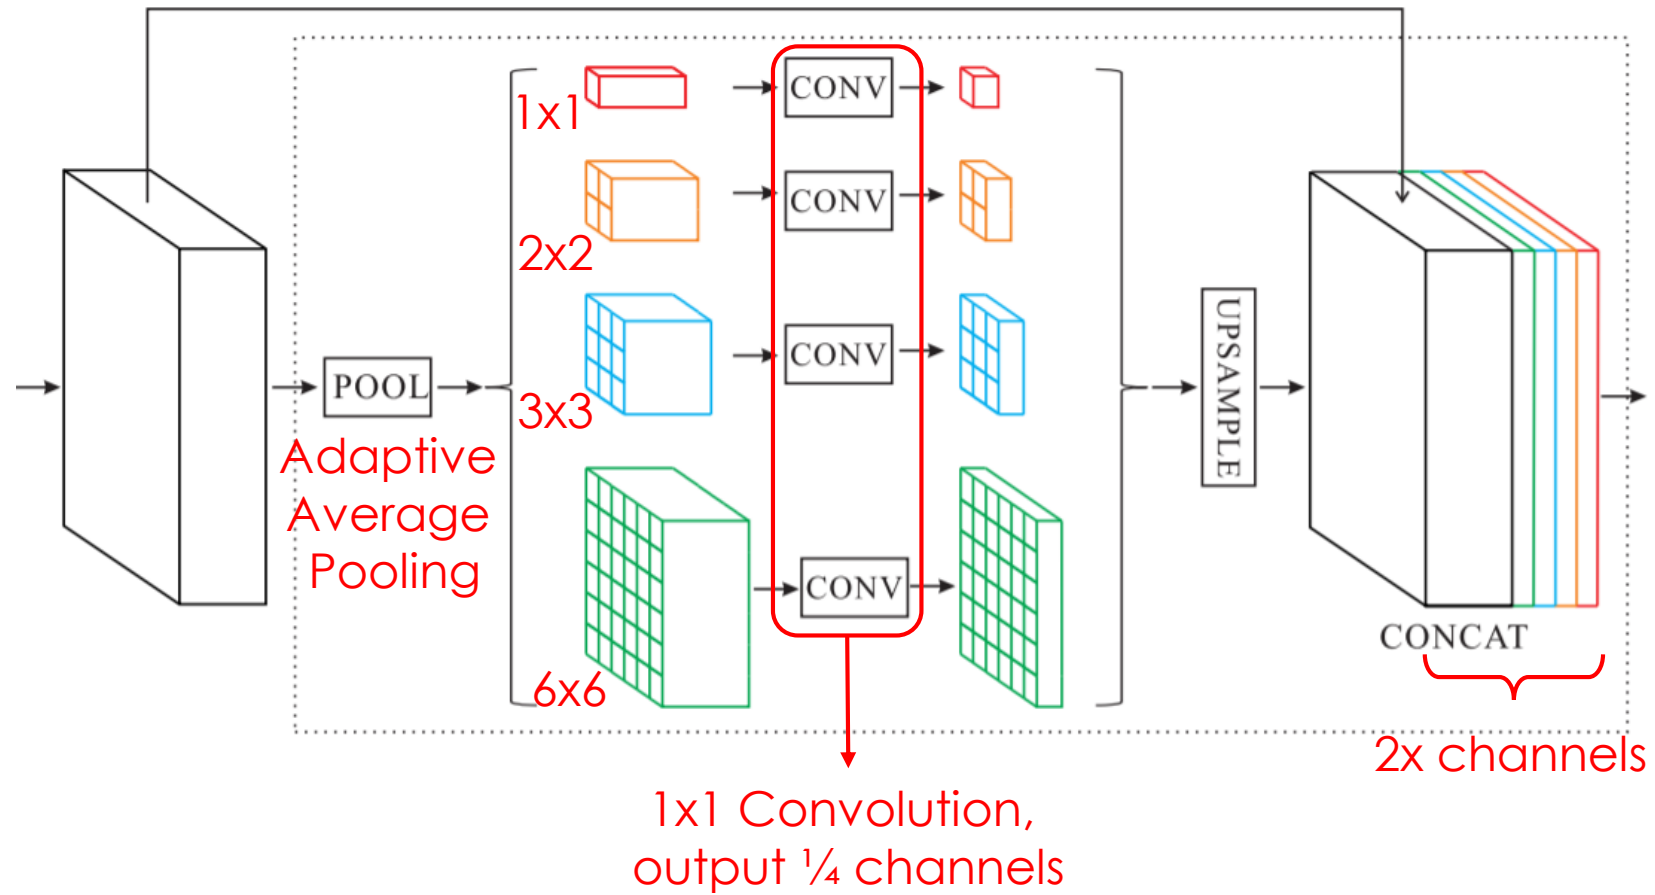
[Hint] Tensor Concatenation
`tc = torch.cat([t1,t2,...],dim)`



Practice - PSPNet



Pyramid Pooling Module



Pyramid Pooling Module

```
# Pyramid Pooling Moudule
self.ppm_size = (16,16)
self.ppm_channel = 512
self.ppm_psize = [1,2,3,6]

self.ppm_pool, self.ppm_conv, self.ppm_up = [], [], []
for psize in self.ppm_psize:
    self.ppm_pool.append(nn.AdaptiveAvgPool2d((psize,psize)))
    self.ppm_conv.append(nn.Conv2d(int(self.ppm_channel), int(self.ppm_channel/len(self.ppm_psize)), kernel_size=1))
    self.ppm_up.append(nn.Upsample(size=self.ppm_size, mode='bilinear', align_corners=True))

self.ppm_pool = nn.ModuleList(self.ppm_pool)
self.ppm_conv = nn.ModuleList(self.ppm_conv)
self.ppm_up = nn.ModuleList(self.ppm_up)
```

```
# Forward
ppm_list = [pre_output[4]]
for i in range(len(self.ppm_psize)):
    output = self.ppm_pool[i](pre_output[4])
    output = self.ppm_conv[i](output)
    output = self.ppm_up[i](self.relu(output))
    ppm_list.append(output)
output = torch.cat(ppm_list,1)
```

Select Your Model

```
[ ] seg_model = EncoderDecoder(pretrained_net=vgg_model, n_class=num_class)
    #seg_model = FCN(pretrained_net=vgg_model, n_class=num_class)
    #seg_model = UNet(pretrained_net=vgg_model, n_class=num_class)
    #seg_model = PSPNet(pretrained_net=vgg_model, n_class=num_class)

seg_model = seg_model.cuda()
criterion = nn.BCEWithLogitsLoss()
optimizer = optim.Adam(seg_model.parameters(), lr=lr)
```

Training and Validation

```
# perform training and validation  
train()
```

```
... epoch: 0, iter: 0, loss: 0.6840  
epoch: 0, iter:10, loss: 0.5426  
epoch: 0, iter:20, loss: 0.4711  
Finish epoch: 0, time elapsed: 73.9815  
pix_acc: 0.6223, meanIoU: 0.2010  
=====
```

epoch	iter	loss
1	0	0.5281
1	10	0.1586
1	20	0.1838

```
Finish epoch: 1, time elapsed: 21.9091  
pix_acc: 0.6395, meanIoU: 0.1937  
=====
```

epoch	iter	loss
2	0	0.1554
2	10	0.1228
2	20	0.1209

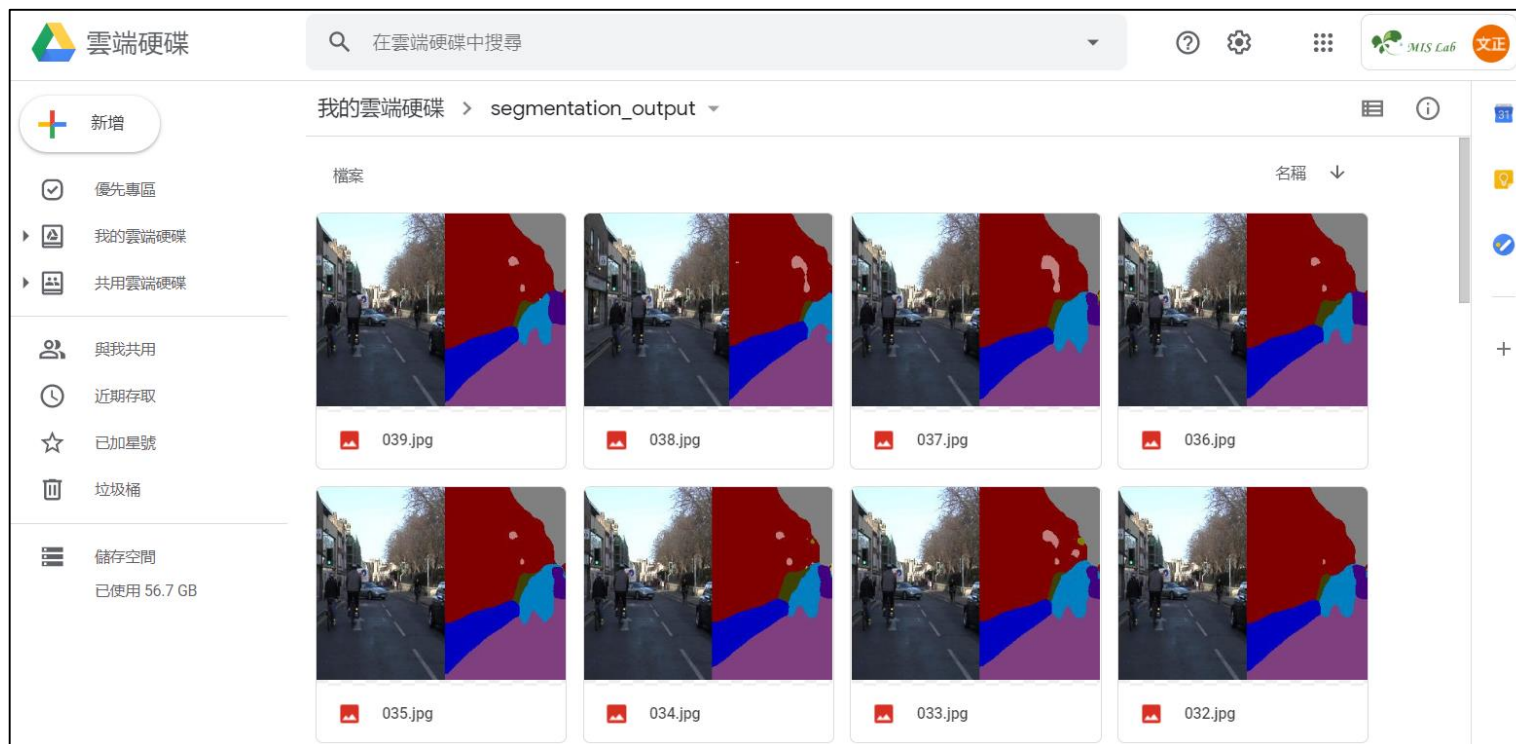
```
Finish epoch: 2, time elapsed: 22.0490  
pix_acc: 0.7600, meanIoU: 0.3024  
=====
```

epoch	iter	loss
3	0	0.1339
3	10	0.1036
3	20	0.1152

```
Finish epoch: 3, time elapsed: 21.8800  
pix_acc: 0.7946, meanIoU: 0.3423  
=====
```

epoch	iter	loss
4	0	0.1201

Output Images:




Experimental Results


Encoder-Decoder



U-Net





FCN



PSPNet



- | | |
|---|------------|
|  | Sky |
|  | Building |
|  | Pole |
|  | Road |
|  | Pavement |
|  | Tree |
|  | SignSymbol |
|  | Fence |
|  | Car |
|  | Pedestrian |
|  | Bicyclist |
|  | Unlabelled |

Q&A



??

