

# Robotic Navigation and Exploration

Week 13: Reinforcement Learning (II)

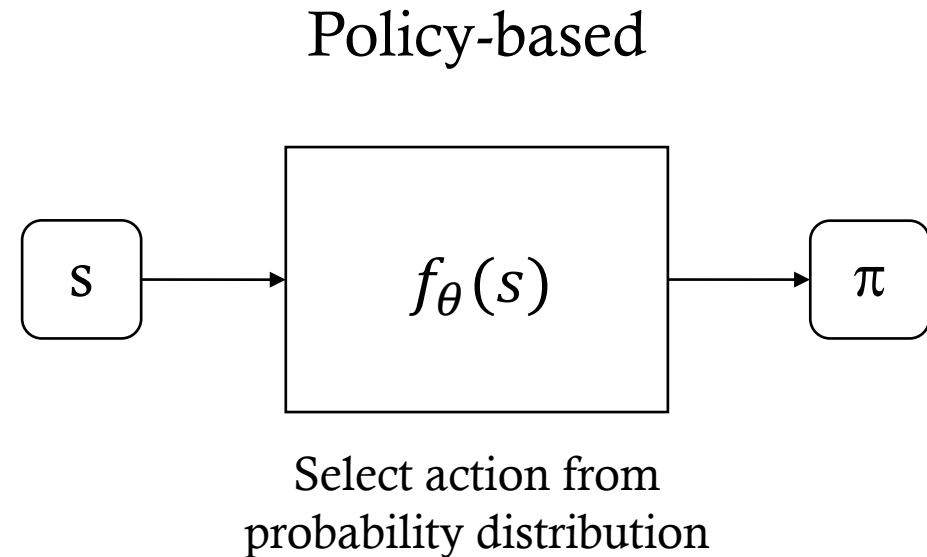
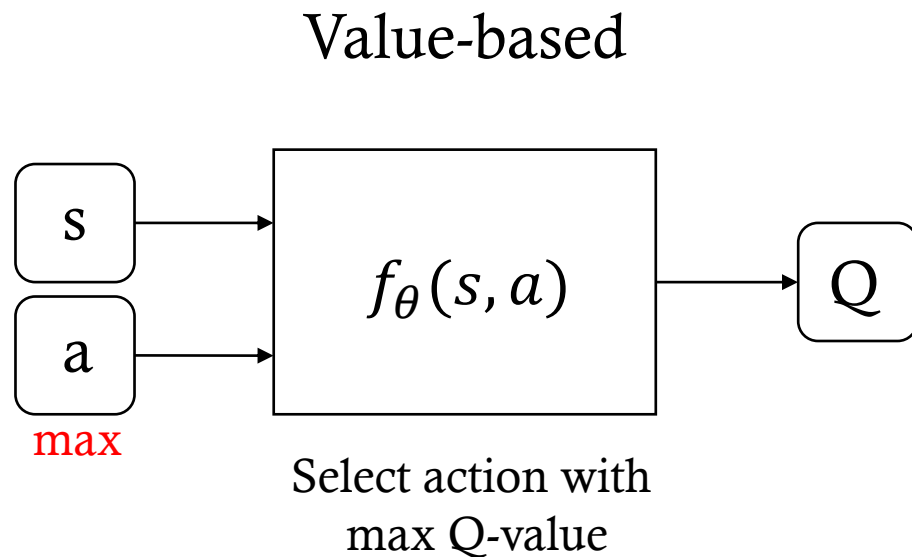
Min-Chun Hu [anitahu@cs.nthu.edu.tw](mailto:anitahu@cs.nthu.edu.tw)  
CS, NTHU

# Outline

- Policy-based Reinforcement Learning
  - Policy Gradient and Actor Critic
  - Sampling Efficiency Problem (TRPO & PPO)
  - Deep Deterministic Policy Gradient (DDPG) & Soft Actor-Critic (SAC)
- Lab6: Model free RL for Mapless Navigation

# Policy-based vs. Value-based Method

- Value-based method aims to learn the action-value function and policy-based method aims to learn the policy function.




# Policy Gradient

- For policy-based method, we construct a function approximator to learn the mapping from state to the action.
- If the computation of gradient for the parameters is achievable, we can utilize iterative optimization method to optimize the policy.

- Definition

- Trajectory  $\tau = \{s_0, a_0, s_1, a_1, \dots, s_T, a_T\}$
- Return  $G(\tau) = r_0 + \gamma r_1 + \gamma^2 r_2 + \dots + \gamma^T r_t$
- Probability of a trajectory

$$P(\tau; \pi) = P(s_0) \prod_{t=0}^{T-1} P(s_{t+1} | s_t, a_t) \pi_{\theta}(a_t | s_t)$$



# Policy Gradient

- Gradient of the policy

$$\begin{aligned}\nabla_{\theta} J(\theta) &= \nabla_{\theta} E \left[ \sum_{t=0}^T \gamma^t r_t \right] \\ &= \nabla_{\theta} \sum_{\tau} P(\tau; \pi_{\theta}(a|s)) * G(\tau) \\ &= \sum_{\tau} \nabla_{\theta} P(\tau; \pi_{\theta}) * G(\tau) \\ &= \sum_{\tau} P(\tau; \pi_{\theta}) \frac{\nabla_{\theta} P(\tau; \pi_{\theta})}{P(\tau; \pi_{\theta})} G(\tau) \\ &= \sum_{\tau} P(\tau; \pi_{\theta}) \nabla_{\theta} \log P(\tau; \pi_{\theta}) * G(\tau) \\ &= \mathbb{E}_{\tau} [\nabla_{\theta} \log P(\tau; \pi_{\theta}) * G(\tau)]\end{aligned}$$

- Probability of a trajectory

$$\begin{aligned}\nabla_{\theta} \log P(\tau; \pi_{\theta}) &= \nabla_{\theta} \log P(s_0) \prod_{t=0}^{T-1} P(s_{t+1}|s_t, a_t) \pi_{\theta}(a_t|s_t) \\ &= \nabla_{\theta} \sum_{t=0}^{T-1} \log P(s_{t+1}|s_t, a_t) + \nabla_{\theta} \sum_{t=0}^{T-1} \log \pi_{\theta}(a_t|s_t) \\ &= \nabla_{\theta} \sum_{t=0}^{T-1} \log \pi_{\theta}(a_t|s_t)\end{aligned}$$

Policy Gradient:

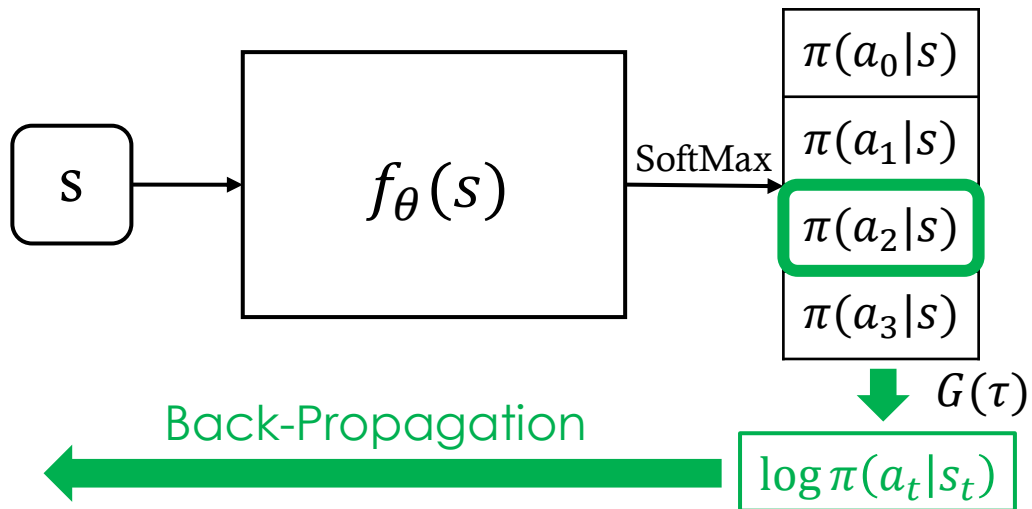
$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau} \left[ \nabla_{\theta} \sum_{t=0}^{T-1} \log \pi_{\theta}(a_t|s_t) * G(\tau) \right]$$

# Policy Gradient

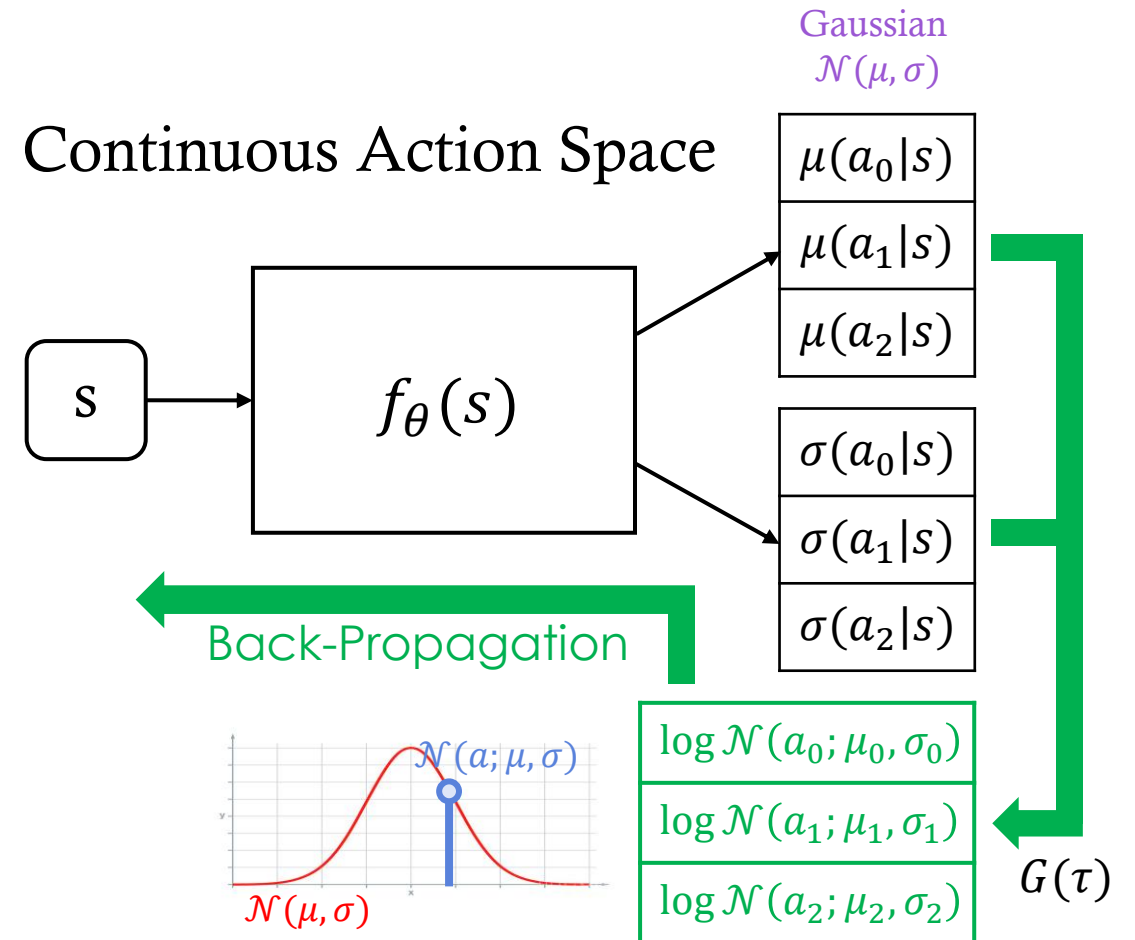
Policy Gradient:

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau} \left[ \nabla_{\theta} \sum_{t=0}^{T-1} \log \pi_{\theta}(\mathbf{a}_t | \mathbf{s}_t) * G(\tau) \right]$$

## Discrete Action Space





## Continuous Action Space



# REINFORCE Algorithm

- Monte-Carlo estimation of the policy gradient



$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau} \left[ \nabla_{\theta} \sum_{t=0}^{T-1} \log \pi_{\theta}(a_t|s_t) * G_t(\tau) \right] \approx \frac{1}{m} \sum_{i=1}^m \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t^i|s_t^i) * G_t(\tau^i)$$

(Episode Update)

# Baseline for Policy Gradient

- Adding an appropriate baseline function can reduce the variance of the estimation for policy gradient.

$$\begin{aligned}\nabla_{\theta} J(\theta) &= \mathbb{E}_{\tau}[\nabla_{\theta} \log \pi_{\theta}(a_t|s_t) * G_t(\tau)] = \mathbb{E}_{\tau}[\nabla_{\theta} \log \pi_{\theta}(a_t|s_t) * (G_t(\tau) - \underbrace{b(s_t)}_{\text{Baseline}})] \\ \Rightarrow \mathbb{E}_{\tau}[\nabla_{\theta} \log \pi_{\theta}(a_t|s_t) * b(s_t)] &= 0\end{aligned}$$

**Proof:**

$$\begin{aligned}&\mathbb{E}_{\tau}[\nabla_{\theta} \log \pi_{\theta}(a_t|s_t) * b(s_t)] \\&= \mathbb{E}_{s_{0:t}, a_{0:t}}[\mathbb{E}_{s_{t+1:T}, a_{t+1:T}}[\nabla_{\theta} \log \pi_{\theta}(a_t|s_t) * b(s_t)]] \\&= \mathbb{E}_{s_{0:t}, a_{0:t}}[b(s_t) * \mathbb{E}_{s_{t+1:T}, a_{t+1:T}}[\nabla_{\theta} \log \pi_{\theta}(a_t|s_t)]] \\&= \mathbb{E}_{s_{0:t}, a_{0:t}}\left[b(s_t) * \sum_{a_t \in A} \sum_{s_{t+1} \in S} \dots \sum_{s_T \in S} \pi_{\theta}(a_t, s_t) P(s_{t+1}|s_t, a_t) \dots P(s_T|s_{T-1}, a_{T-1}) (\nabla_{\theta} \log \pi_{\theta}(a_t|s_t))\right] \\&= \mathbb{E}_{s_{0:t}, a_{0:t}}\left[b(s_t) * \sum_{a_t \in A} \pi_{\theta}(a_t|s_t) \nabla_{\theta} \log \pi_{\theta}(a_t|s_t) \sum_{s_{t+1} \in S} P(s_{t+1}|s_t, a_t) \sum_{a_{t+1} \in A} \dots \sum_{s_T \in S} P(s_T|s_{T-1}, a_{T-1})\right] \\&= \mathbb{E}_{s_{0:t}, a_{0:t}}[b(s_t) * \mathbb{E}_{a_t}[\nabla_{\theta} \log \pi_{\theta}(a_t|s_t)]] = \mathbb{E}_{s_{0:t}, a_{0:t}}[b(s_t) * 0] = 0\end{aligned}$$



$$\begin{aligned} & \mathbb{E}_{a_t}[\nabla_{\theta} \log \pi_{\theta}(a_t|s_t)] \\ &= \sum_{a_t \in A} \pi_{\theta}(a_t|s_t) \nabla_{\theta} \log \pi_{\theta}(a_t|s_t) \\ &= \sum_{a_t \in A} \pi_{\theta}(a_t|s_t) \frac{\nabla_{\theta} \pi_{\theta}(a_t|s_t)}{\pi_{\theta}(a_t|s_t)} \\ &= \sum_{a_t \in A} \nabla_{\theta} \pi_{\theta}(a_t|s_t) \\ &= \nabla_{\theta} \sum_{a_t \in A} \pi_{\theta}(a_t|s_t) \\ &= \nabla_{\theta} 1 \\ &= 0 \end{aligned}$$

# Baseline for Policy Gradient

- Consider the variance of policy gradient

$$\begin{aligned} & \text{Var} \left( \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) * (G_t(\tau) - b(s_t)) \right) \\ & \approx \sum_{t=0}^{T-1} \mathbb{E}_{\tau} \left[ \left( \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) * (G_t(\tau) - b(s_t)) \right)^2 \right] \\ & \approx \sum_{t=0}^{T-1} \mathbb{E}_{\tau} [(\nabla_{\theta} \log \pi_{\theta}(a_t | s_t))^2] * \mathbb{E}_{\tau} \left[ \left( (G_t(\tau) - b(s_t)) \right)^2 \right] \end{aligned}$$

Setting  $b(s_t) \approx \mathbb{E}_{\tau}[G_t(\tau)]$  to approximate the expected return will have low variance.

# Off-Policy with Importance Sampling

- We compute the policy gradient of  $\pi_\theta$  by collecting the data under the probability of policy  $\pi_\theta$ . Once we update the policy, we need to collect the data again.

$$\nabla_\theta J(\theta) = \mathbb{E}_{\tau \sim p_\theta(\tau)} [G(\tau) \nabla_\theta \log p_\theta(\tau)]$$

- If we want to reuse the data sampled before updating the policy, we can apply the **importance sampling** technique to estimate the policy gradient.


$$\nabla_\theta J(\theta) = \mathbb{E}_{\tau \sim p_{\theta'}(\tau)} \left[ \frac{p_\theta(\tau)}{p_{\theta'}(\tau)} G(\tau) \nabla_\theta \log p_\theta(\tau) \right]$$

Sample from  $p_{\theta'}(\tau)$       Importance Weight


$$\mathbb{E}_{x \sim p}[f(x)] = \mathbb{E}_{x \sim q} \left[ \frac{p(x)}{q(x)} f(x) \right]$$

# Importance Sampling Review

- Important sampling adopts discrete multinomial to approximate arbitrary distribution. More sampling particles will have more accurate approximation.



1. Sampling  $x_i$  from  $\pi(x)$
2. Calculate  $w_i = \frac{p(x_i)}{\pi(x_i)}$
3. Sampling  $x$  from  $\text{mul}(x_i, w_i)$



# Off-Policy with Importance Sampling

$$\begin{aligned}\nabla_{\theta} J(\theta) &= \mathbb{E}_{\tau \sim p_{\theta}(\tau)} [G(\tau) \nabla_{\theta} \log p_{\theta}(\tau)] \\ &= \mathbb{E}_{\tau \sim p_{\theta'}(\tau)} \left[ \frac{p_{\theta}(\tau)}{p_{\theta'}(\tau)} G^{\theta'}(\tau) \nabla_{\theta} \log p_{\theta}(\tau) \right] \\ &= \mathbb{E}_{\tau \sim p_{\theta'}(\tau)} \left[ \frac{p_{\theta}(s_t, a_t)}{p_{\theta'}(s_t, a_t)} G^{\theta'}(\tau) \nabla_{\theta} \log p_{\theta}(s_t, a_t) \right] \\ &= \mathbb{E}_{\tau \sim p_{\theta'}(\tau)} \left[ \frac{p_{\theta}(a_t | s_t)}{p_{\theta'}(a_t | s_t)} \frac{\cancel{p_{\theta}(s_t)}}{\cancel{p_{\theta'}(s_t)}} G^{\theta'}(\tau) \nabla_{\theta} \log p_{\theta}(s_t, a_t) \right]\end{aligned}$$

$$J^{\theta'}(\theta) = \mathbb{E}_{\tau \sim p_{\theta'}(\tau)} \left[ \frac{p_{\theta}(a_t | s_t)}{p_{\theta'}(a_t | s_t)} G^{\theta'}(\tau) \right]$$

Trust Region Policy Optimization (TRPO)

$$J_{TRPO}^{\theta'}(\theta) = \mathbb{E}_{\tau \sim p_{\theta'}(\tau)} \left[ \frac{p_{\theta}(a_t | s_t)}{p_{\theta'}(a_t | s_t)} G^{\theta'}(\tau) \right]$$

Subject to  $KL(\theta, \theta') < \delta$

Proximal Policy Optimization (PPO)

$$\nabla_{\theta'}^{PPO} J(\theta) = J^{\theta'}(\theta) - \beta KL(\theta, \theta')$$

$$J^{\theta'}(\theta) = \mathbb{E}_{\tau \sim p_{\theta'}(\tau)} \left[ \frac{p_{\theta}(a_t | s_t)}{p_{\theta'}(a_t | s_t)} G^{\theta'}(\tau) \right]$$

# Proximal Policy Optimization



Proximal Policy Optimization (PPO)

$$\nabla_{\theta'}^{PPO} J(\theta) = J^{\theta'}(\theta) - \beta KL(\theta, \theta')$$

$$J^{\theta'}(\theta) = \mathbb{E}_{\tau \sim p_{\theta'}(\tau)} \left[ \frac{p_{\theta}(a_t|s_t)}{p_{\theta'}(a_t|s_t)} G^{\theta'}(\tau) \right]$$

Proximal Policy Optimization (Clip Version)

$$\nabla_{\theta'}^{PPO2} J(\theta) \approx \sum_{s_t, a_t} \min \left( \frac{p_{\theta}(a_t|s_t)}{p_{\theta'}(a_t|s_t)} G^{\theta'}(\tau), \text{clip} \left( \frac{p_{\theta}(a_t|s_t)}{p_{\theta'}(a_t|s_t)}, 1 - \epsilon, 1 + \epsilon \right) G^{\theta'}(\tau) \right)$$



# Actor-Critic


$$\nabla_{\theta} J(\theta) = E_{\pi} \left[ \nabla_{\theta} \sum_{t=0}^{T-1} \log \pi_{\theta}(a_t | s_t) G_t(\tau) \right] \quad \text{Policy Gradient}$$

$$= E_{\pi} \left[ \nabla_{\theta} \sum_{t=0}^{T-1} \log \pi_{\theta}(a_t | s_t) Q(s_t, a_t) \right] \quad \text{Actor-Critic}$$


$$= E_{\pi} \left[ \nabla_{\theta} \sum_{t=0}^{T-1} \log \pi_{\theta}(a_t | s_t) [Q(s_t, a_t) - B(s_t)] \right] \quad \text{Actor-Critic with Baseline}$$

$$= E_{\pi} \left[ \nabla_{\theta} \sum_{t=0}^{T-1} \log \pi_{\theta}(a_t | s_t) [Q(s_t, a_t) - V(s_t)] \right] \quad \text{Advantage Actor-Critic}$$


$$= E_{\pi} \left[ \nabla_{\theta} \sum_{t=0}^{T-1} \log \pi_{\theta}(a_t | s_t) [r_t + \gamma V(s_{t+1}) - V(s_t)] \right] \quad \text{TD Actor-Critic}$$



Policy Iteration



# TD Actor-Critic



*Actor (Policy) Loss:*

$$L(\theta) =$$

$$\mathbb{E}[-\log \pi_{\theta}(a|s) (\textcolor{red}{r + \gamma V(s') - V(s)}) - \textcolor{blue}{\beta H(\pi_{\theta}(a|s))}]$$

Advantage Entropy Term  
(Encourage Exploration)



*Critic (Value) Loss:*

$$L(\phi) = \mathbb{E}[\frac{1}{2} (\textcolor{red}{r + \gamma V(s') - V(s)})^2]$$

TD-Error




# Deterministic Policy Gradient (DPG)

- The DPG algorithm maintains an actor function  $\mu(\mathbf{s} | \boldsymbol{\theta})$  which specifies the current policy by deterministically mapping states to a specific action.

- **Stochastic Policy**

$$Q^\pi(s_t, a_t) = \mathbb{E}_{r_t, s_{t+1} \sim P}[r(s_t, a_t) + \gamma \mathbb{E}_{a_{t+1} \sim \pi}[Q^\pi(s_{t+1}, a_{t+1})]]$$

- **Deterministic Policy**

$$Q^\mu(s_t, a_t) = \mathbb{E}_{r_t, s_{t+1} \sim P}[r(s_t, a_t) + \gamma Q^\mu(s_{t+1}, \mu(s_{t+1}))]$$


- Without the stochastic property of the policy distribution, the policy gradient can be estimated simpler by chain rules.

$$\begin{aligned} \nabla_{\boldsymbol{\theta}} J(\mu_{\boldsymbol{\theta}}) &\approx \int_s P^\mu(s) \nabla_{\boldsymbol{\theta}} \mu_{\boldsymbol{\theta}}(s) \nabla_a Q^\mu(s, a) \Big|_{a=\mu_{\boldsymbol{\theta}}(s)} ds \\ &= \mathbb{E}_{s \sim P^\mu} [\nabla_{\boldsymbol{\theta}} \mu_{\boldsymbol{\theta}}(s) \nabla_a Q^\mu(s, a) |_{a=\mu_{\boldsymbol{\theta}}(s)}] \end{aligned}$$

# On-policy and Off-policy DPG

- On-policy DPG

$$\begin{aligned}\nabla_{\theta} J(\mu_{\theta}) &= \int_s P^{\mu}(s) \nabla_{\theta} \mu_{\theta}(s) \nabla_a Q^{\mu}(s, a) ds \\ &= \mathbb{E}_{s \sim P^{\mu}} [\nabla_{\theta} \mu_{\theta}(s) \nabla_a Q^{\mu}(s, a) |_{a=\mu_{\theta}(s)}]\end{aligned}$$

- Off-policy DPG

$$\begin{aligned}\nabla_{\theta} J_{\beta}(\mu_{\theta}) &\approx \int_s P^{\beta}(s) \nabla_{\theta} \mu_{\theta}(s) \nabla_a Q^{\mu}(s, a) ds \\ &= \mathbb{E}_{s \sim P^{\beta}} [\nabla_{\theta} \mu_{\theta}(s) \nabla_a Q^{\mu}(s, a) |_{a=\mu_{\theta}(s)}]\end{aligned}$$

## Update Parameters

$$Q^w(s, a) \approx Q^{\mu}(s, a)$$

$$\delta_t = r_t + \gamma Q^w(s_{t+1}, a_{t+1}) - Q^w(s_t, a_t)$$

$$w_{t+1} = w_t + \alpha_w \delta_t \nabla_w Q^w(s_t, a_t)$$

$$\theta_{t+1} = \theta_t + \alpha_{\theta} \nabla_{\theta} \mu_{\theta}(s_t) \nabla_a Q^w(s_t, a_t) |_{a=\mu_{\theta}(s)}$$

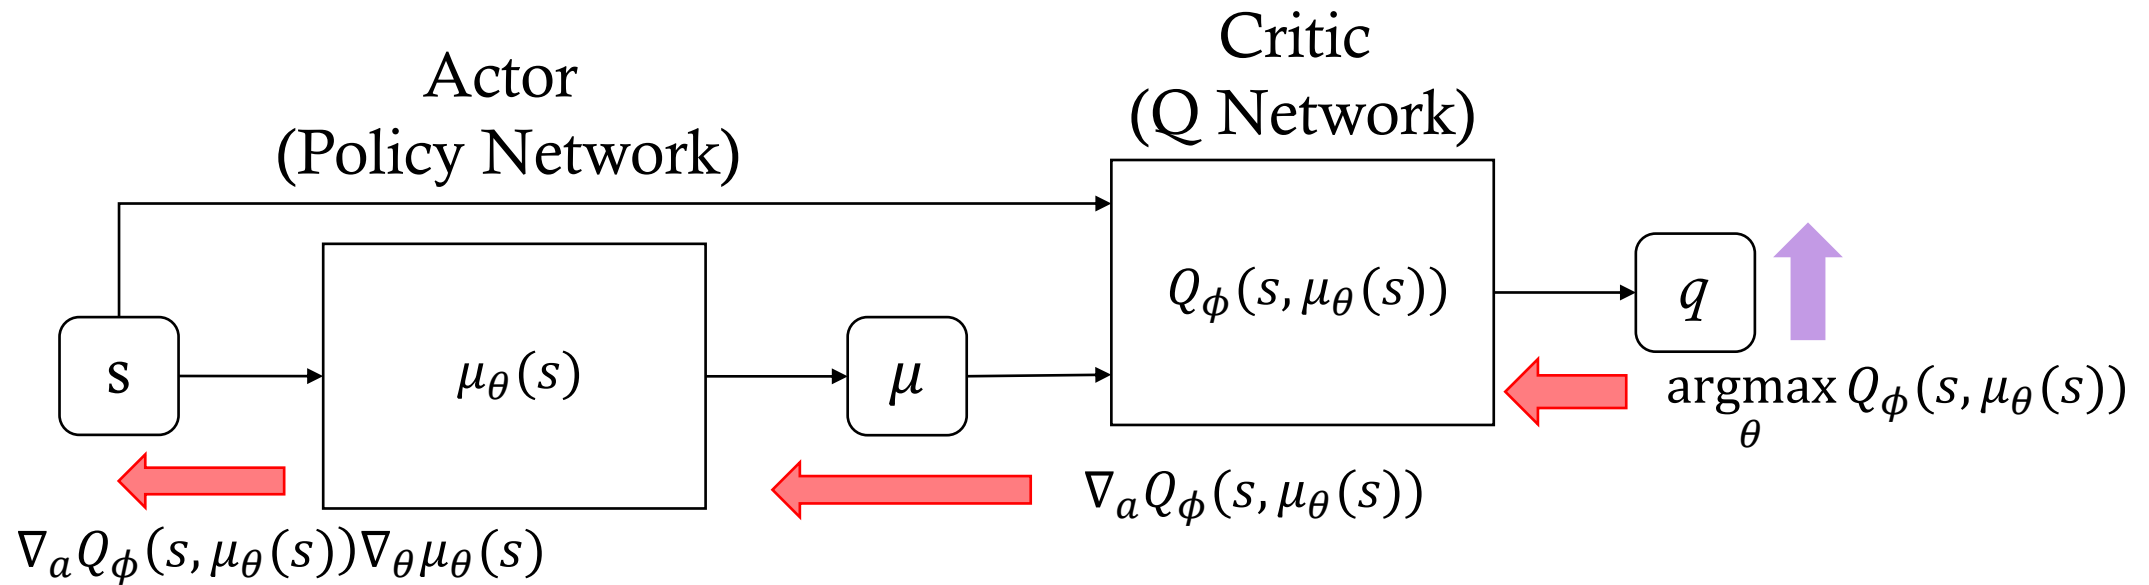
Assume  $\underset{a}{\operatorname{argmax}} Q(s, a) = \mu_{\theta}(s_{t+1})$

$$\delta_t = r_t + \boxed{\gamma Q^w(s_{t+1}, \mu_{\theta}(s_{t+1}))} - Q^w(s_t, a_t)$$

$$w_{t+1} = w_t + \alpha_w \delta_t \nabla_w Q^w(s_t, a_t)$$

$$\theta_{t+1} = \theta_t + \alpha_{\theta} \nabla_{\theta} \mu_{\theta}(s_t) \nabla_a Q^w(s_t, a_t) |_{a=\mu_{\theta}(s)}$$

# Deterministic Policy Gradient (DPG)

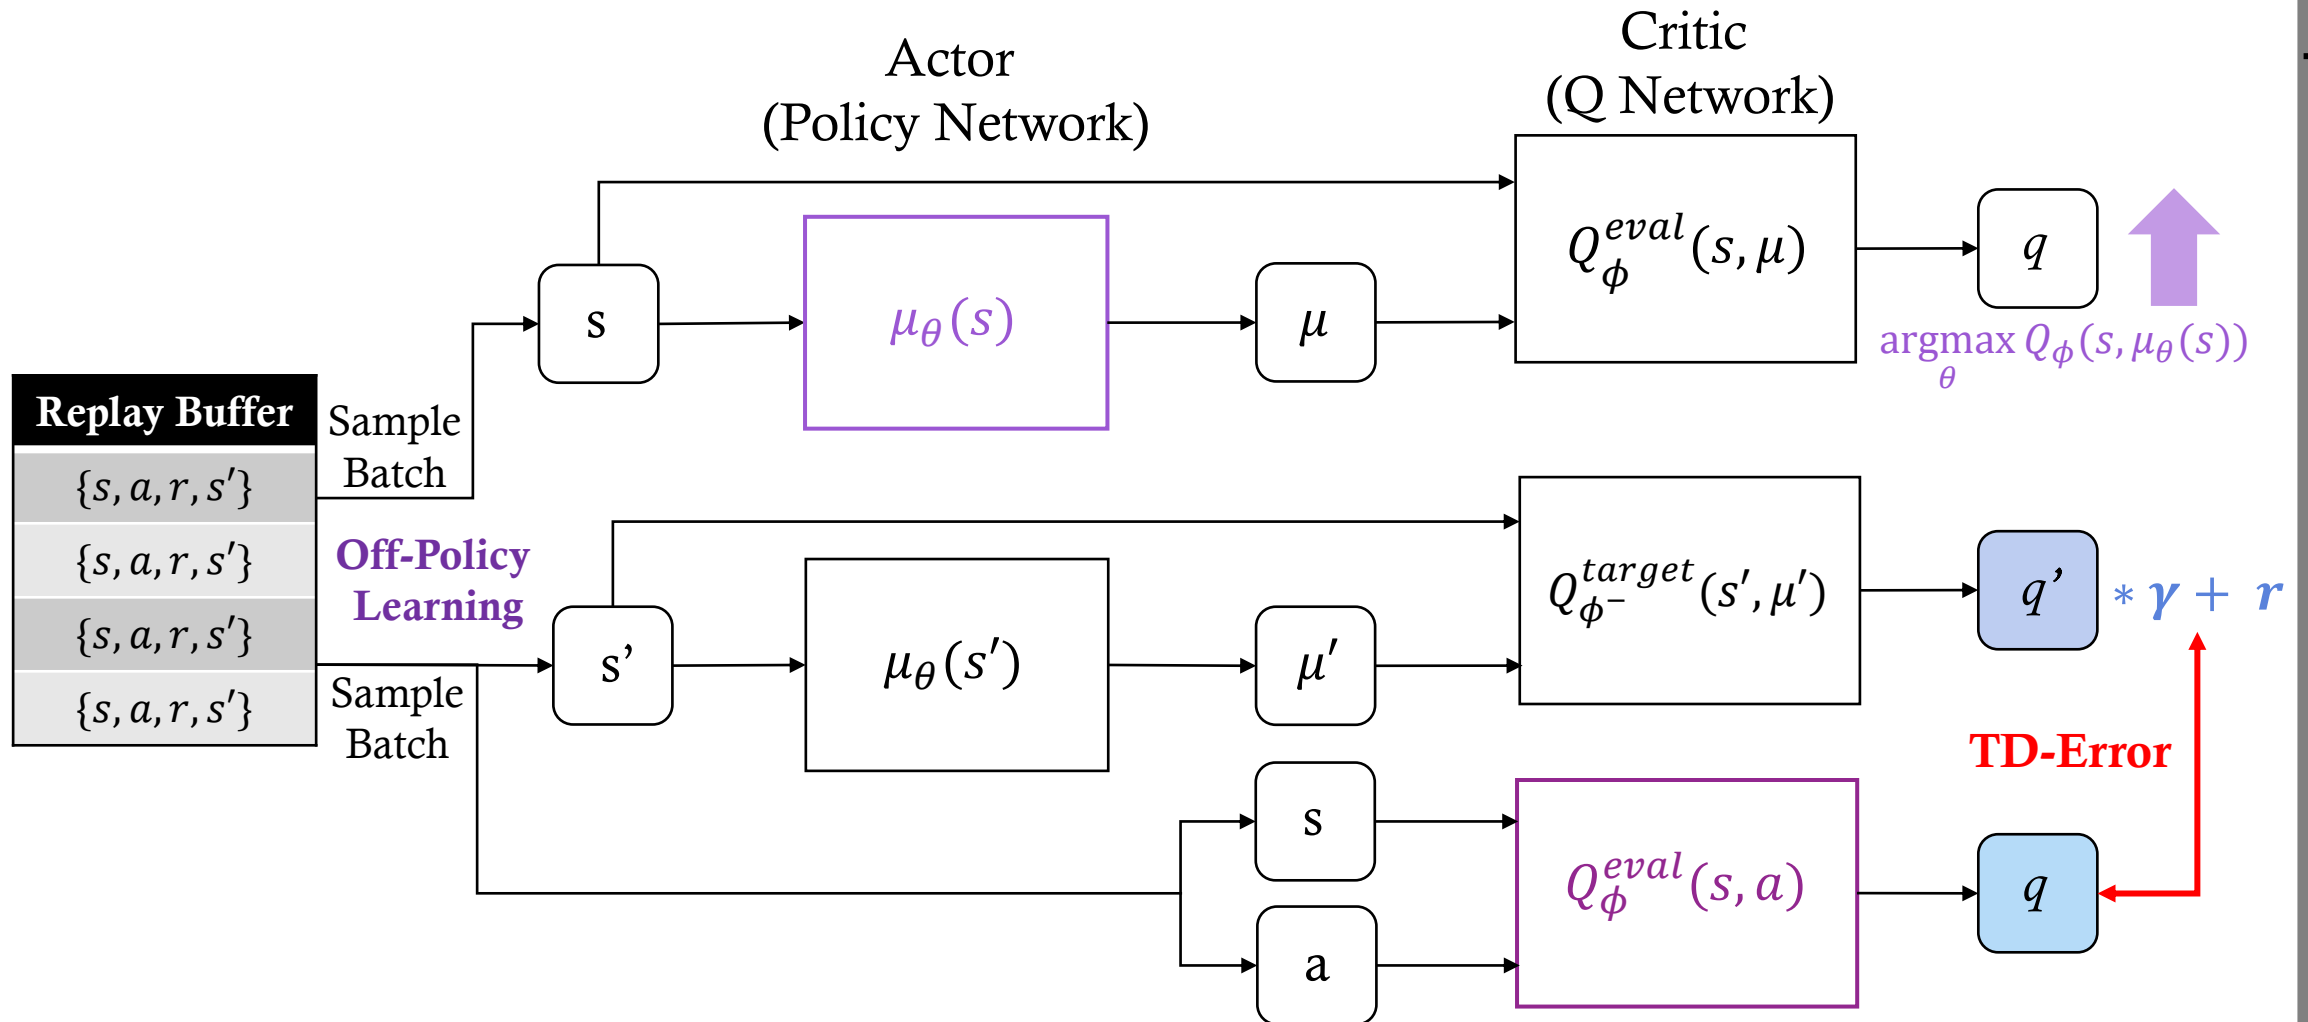


*Actor (Policy) Loss:*  
$$L(\theta) = \mathbb{E}[-Q_\phi(s, \mu_\theta(s))]$$



*Critic (Q) Loss:*  
$$L(\phi) = \mathbb{E}\left[\frac{1}{2} \left( r + \gamma Q_\phi(s', \mu(s')) - Q(s, a) \right)^2 \right]$$
  
**TD-Error**

**Off-Policy  
Learning**

# Deep Deterministic Policy Gradient (DDPG)



# Exploration Issue



- The formulation of classic RL framework implies an optimal deterministic policy, while it conflict to the stochastic policy that encourages the exploration of the environment.
- We usually add an regularized entropy term to maintain the stochastic property in stochastic policy algorithm such as policy gradient or actor-critic.

$$L_{actor}(\theta) = \mathbb{E}[-\log \pi_{\theta}(a|s) A(s) - \beta H(\pi_{\theta}(a|s))]$$

- An extension of classic RL framework is **Maximum Entropy Reinforcement Learning Framework**, which take the long-term entropy of the policy distribution into consideration to ensure the exploration of the environment.

$$\pi_{MaxEnt}^* = \operatorname{argmax}_{\pi} \sum_t \mathbb{E}_{\pi} [r(s_t, a_t) + \alpha \mathcal{H}(\pi(\cdot | s_t))]$$

Consider the future entropy

# Maximum Entropy RL Framework

$$\pi_{MaxEnt}^* = \operatorname{argmax}_{\pi} \sum_t \mathbb{E}_{\pi} [r(s_t, a_t) + \alpha \mathcal{H}(\pi(\cdot | s_t))]$$

Maximum Entropy  
Objective

$$V_{soft}(s_t) = \mathbb{E}_{\pi} [Q_{soft}(s_t, a_t) - \alpha \log \pi(a_t | s_t)]$$

Energy-based  
Policy

$$\begin{aligned} \varepsilon(s_t, a_t) &= -\frac{1}{\alpha} Q_{soft}(s_t, a_t) \\ \pi(a_t | s_t) &\propto \exp(Q_{soft}(s_t, a_t)) \end{aligned}$$

Soft Value  
Function

$$\begin{aligned} V_{soft}(s_t) &= \log \int \exp\left(\frac{1}{\alpha} Q_{soft}(s_t, a)\right) da \\ &= \operatorname{soft max}_a Q_{soft}(s_t, a) \end{aligned}$$

# Soft Actor-Critic

- Soft Actor Critic (SAC) algorithm is an off-policy algorithm based on the soft policy iteration of maximum entropy RL framework.
- SAC can be seen as the stochastic extension of the DDPG, which adopt the **reparameterization trick** to estimate the gradient of gaussian policy.

*Actor (Policy) Loss:*

$$L(\theta) = \mathbb{E}[\alpha \log \pi_{\theta}(f_{\theta}(\epsilon_t; s_t) | s_t) - Q_{\phi}(s_t, f_{\theta}(\epsilon_t; s_t))]$$

$$a_t = f_{\theta}(\epsilon_t; s_t) \text{ (Reparameterize)}$$

*Critic (Q) Loss:*

$$L(\phi) = \mathbb{E} \left[ \frac{1}{2} \left( -\gamma \alpha \log \pi(\cdot | s_{t+1}) + r(s_t, a_t) + \gamma Q_{\phi}(s_{t+1}, f_{\theta}(\epsilon_{t+1}; s_{t+1})) - Q_{\phi}(s_t, a_t) \right)^2 \right]$$

*Actor Loss (DDPG):*



$$L(\theta) = \mathbb{E}[-Q_{\phi}(s, \mu_{\theta}(s))]$$

*Critic Loss (DDPG):*


$$L(\phi) = \mathbb{E} \left[ \frac{1}{2} \left( r + \gamma Q_{\phi}(s', \mu(s')) - Q(s, a) \right)^2 \right]$$

# Reparameterization


- Unbias gradient estimator for Gaussian sampling



Original

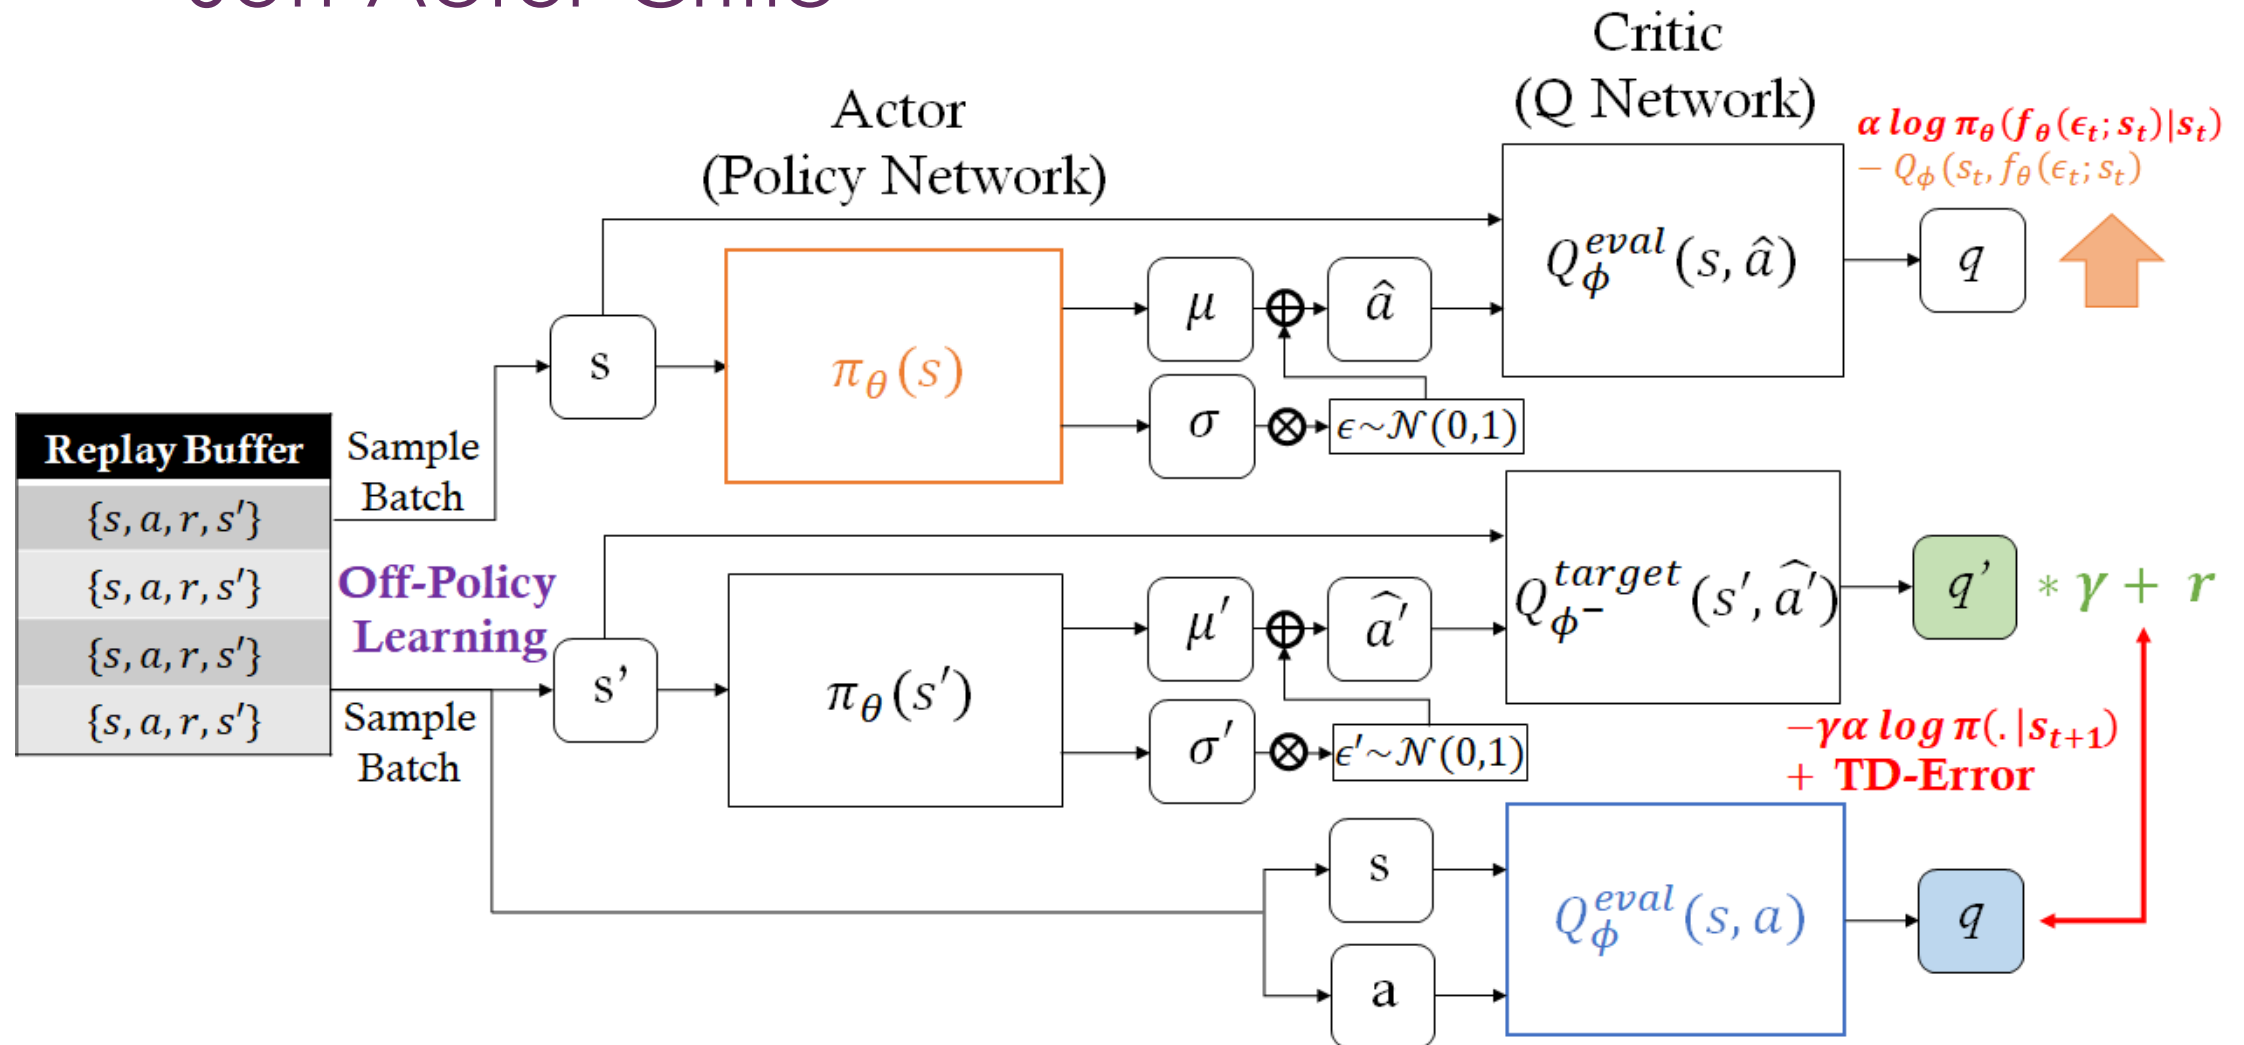


Reparameterize





# Soft Actor-Critic



# Reinforcement Learning Algorithms

